

# MODEL DOCUMENTATION

## Multi-Layer Perceptron (MLP)

### 1. Model Used

The model used in this project is a **Feedforward Neural Network (FNN)**, implemented using scikit-learn's `MLPClassifier`. It is a type of **Multi-Layer Perceptron (MLP)**, which is a supervised deep learning algorithm consisting of an input layer, one or more hidden layers, and an output layer.

### 2. Why Feedforward Neural Network (FNN)?

**Captures Nonlinear Relationships:** Able to model complex, nonlinear interactions among features rather than relying on simple assumptions.

**Representation Learning:** Learns hierarchical feature representations from TF-IDF vectors, which enhances prediction quality.

**Flexible Architecture:** The number of layers and neurons can be customized to balance performance with computational cost.

**Generalization Power:** With techniques like dropout and early stopping, FNNs can adapt well to unseen test data.

**Scalable to Large Data:** More effective when working with large-scale datasets, making it suitable for balanced and imbalanced scenarios.

### 3. Working Principle

The Feedforward Neural Network (FNN) processes data in sequential layers:

1. **Input Layer:** Takes in the TF-IDF features of the text data.
2. **Hidden Layers:** Each neuron applies a weighted sum of inputs followed by a non-linear activation function (e.g., ReLU). This allows the network to capture complex patterns and relationships between words and ratings.
3. **Output Layer:** Uses a softmax activation function to generate probability scores for each rating class (1★ to 5★). The class with the highest probability is assigned as the prediction.

The model learns by adjusting weights through **backpropagation** and optimization (Adam solver in this case) to minimize the loss function.

## 5. Strengths of Feedforward Neural Network (FNN) in This Project

- Learns **nonlinear relationships** between features, going beyond simple assumptions.
- Works effectively with **TF-IDF features** by capturing richer interactions.
- Flexible architecture — hidden layers and neurons can be tuned for better performance.
- Provides **probabilistic outputs** via softmax, which can be interpreted as class confidence.
- Scales well to large datasets (e.g., 25,000 samples per class).

## 6. Limitations

- Requires **longer training time** compared to simpler models.
- Sensitive to hyperparameters (layers, neurons, learning rate, batch size).
- Less interpretable compared to simpler probabilistic models.
- Can **overfit** if not controlled with regularization techniques (e.g., early stopping, dropout).
- Performance depends heavily on consistent preprocessing (same TF-IDF vectorizer).

## 7. Hyperparameter Tuning

In this project, we performed hyperparameter tuning for the **FNN classifier** to balance training efficiency and accuracy.

- **Hyperparameters tuned:**
  - `hidden_layer_sizes` (number of layers and neurons per layer)
  - `activation` (ReLU chosen for efficiency and nonlinearity)
  - `batch_size` (controls training updates per step)
  - `max_iter` (maximum epochs to train)
  - `early_stopping=True` (to prevent overfitting)
- **Method used:** Manual experimentation with multiple architectures.
- **Best configuration found:**
  - Hidden layers: (256, 128)

- Activation: ReLU
- Batch size: 128
- Early stopping enabled
- **Impact:** This configuration achieved around **61% accuracy** on the balanced dataset, significantly improving class-wise prediction consistency compared to simpler baselines.

## 8. Conclusion

In this project, a **Feedforward Neural Network (FNN)** was chosen because it provides a good balance between **accuracy, scalability, and computational efficiency** for text-based rating prediction tasks. While more advanced architectures like RNNs or Transformers could offer higher performance, FNN was selected due to its **simplicity, effectiveness with TF-IDF features, and suitability for large datasets**. All training, evaluation, and inference steps were carried out using FNN as the core deep learning model.

### 1. Balance the Dataset

# Balance the dataset to have equal number of reviews per rating

```
balanced_df = balance_dataset(review, target_col="Rating", n_samples=25000)
```

### 2. Train-Test Split & TF-IDF Vectorization

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
X = balanced_df['Review_text']
```

```
y = balanced_df['Rating']
```

```
# Stratified train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.2, stratify=y, random_state=42
```

```
)
```

```
# TF-IDF Vectorization
```

```
tfidf = TfidfVectorizer(max_features=25000, ngram_range=(1,2), stop_words='english')
```

```
X_train_tfidf = tfidf.fit_transform(X_train)
```

```
X_test_tfidf = tfidf.transform(X_test)
```

### **3. Train Feedforward Neural Network (MLPClassifier)**

```
from sklearn.neural_network import MLPClassifier
```

```
# Initialize and train the model
```

```
mlp_model = MLPClassifier(hidden_layer_sizes=(256,128),
```

```
    activation='relu',
```

```
    solver='adam',
```

```
    batch_size=128,
```

```
    max_iter=20,
```

```
    early_stopping=True,
```

```
    verbose=True,
```

```
    random_state=42)
```

```
mlp_model.fit(X_train_tfidf, y_train)
```

### **4. Evaluate the Model & Hyperparameter Tuning**

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.metrics import accuracy_score, classification_report
```

```
# Hyperparameter tuning for MLP
```

```
param_grid = {
```

```
'hidden_layer_sizes': [(128,), (256,128), (512,256,128)],  
'batch_size': [64, 128],  
'max_iter': [20, 50]  
}
```

```
grid = GridSearchCV(MLPClassifier(activation='relu', solver='adam', early_stopping=True,  
random_state=42),
```

```
    param_grid,  
    cv=3,  
    scoring='accuracy',  
    verbose=2,  
    n_jobs=-1)
```

```
grid.fit(X_train_tfidf, y_train)
```

```
best_mlp_model = grid.best_estimator_  
print("Best Parameters:", grid.best_params_)
```

```
# Evaluate tuned model
```

```
y_pred = best_mlp_model.predict(X_test_tfidf)  
print("FNN Accuracy (tuned):", accuracy_score(y_test, y_pred))  
print(classification_report(y_test, y_pred))
```